

# Notes on and around ISLP

Calvin Khor

Last compiled: October 26, 2023

## 1 Resources

1. The website for the book:
2. Solutions to ISLR:
3. Videos accompanying ISLR:
4. My errata: <https://docs.oracle.com/javase/8/docs/api/java/lang/>
5. These notes: <https://dotty.epfl.ch/3.0.0/api>

## 2 Statistical Learning

### 2.1 Definitions and Notation

$X_1, \dots, X_p$  denotes training data,  $X_0$  is test (out-sample) data. When its capital letters its either a random variable or an abstract model? The  $p$  denotes the number of predictors (i.e. independent variables/features/variables). The model is abstractly

$$Y = f(X) + \epsilon \quad (2.1)$$

$Y$  is the response/dependent variable.  $\epsilon$  is noise inherent to the model: WLOG  $\mathbb{E} \epsilon = 0$ ,  $\text{Var} \epsilon = 1$ , and independent from the predictors (i.e. we cannot predict the error)

We denote the number of observations by  $n$ , so that the observations of the predictors are  $x_j = (x_{1j}, x_{2j}, \dots, x_{nj})$ , for  $j = 1, \dots, p$ .

\*Parametric and non-parametric models\*:

\*Prediction Accuracy\*:

\*Model Interpretability\*:

\*Test MSE\*:

Variance: "the amount by which  $\hat{f}$  would change if we estimated it using a different training data set." - i.e. variance by treating the training data as random variables

Bias: the error that is introduced by approximating a real-life problem, which may be extremely complicated, by a much simpler model.

$$E|y_0 - \hat{f}(x_0)|^2 = \text{Var} \hat{f}(x_0) + (\text{Bias} \hat{f}(x_0))^2 + \text{Var} \epsilon \quad (2.7)$$

—

Important quote from book:

Here the notation  $\mathbb{E} |y_0 - \hat{f}(x_0)|^2$  defines the expected test MSE at  $x_0$ , expected test MSE and refers to the average test MSE that we would obtain if we repeatedly estimated  $f$  using a large number of training sets, and tested each at  $x_0$ . The overall expected test MSE can be computed by averaging  $\mathbb{E} |y_0 - \hat{f}(x_0)|^2$  over all possible values of  $x_0$  in the test set.

Derivation of (2.7):

$$\begin{aligned} \text{LHS}(2.7) &= \mathbb{E} |f(x_0) - \hat{f}(x_0) + \epsilon|^2 \\ &= \mathbb{E} |f(x_0) - \mathbb{E} \hat{f}(x_0) + (\mathbb{E} \hat{f}(x_0) - \hat{f}(x_0)) + \epsilon|^2 \\ &= \mathbb{E} |f(x_0) - \mathbb{E} \hat{f}(x_0)|^2 + \mathbb{E} |\epsilon|^2 + \mathbb{E} |\hat{f}(x_0) - \mathbb{E} \hat{f}(x_0)|^2 + 2 \mathbb{E} \epsilon (f(x_0) - \hat{f}(x_0)) \\ &= |f(x_0) - \mathbb{E} \hat{f}(x_0)|^2 + \mathbb{E} |\epsilon|^2 + \mathbb{E} |\hat{f}(x_0) - \mathbb{E} \hat{f}(x_0)|^2 + 2 \mathbb{E} \epsilon (f(x_0) - \hat{f}(x_0)) \\ &=: \text{Bias}^2 + \text{Var } \epsilon + \text{Var } \hat{f}(x_0) + 0 \text{ (since } \epsilon \text{ is independent)} = \text{RHS}(2.7). \end{aligned}$$

## 2.2 Bias-Variance Trade-off

## 3 Regression Models

4

5

6

7

## 8 Tree-based Methods

9

10

## 11 Python

### 11.1 Classes

<https://stackoverflow.com/questions/100003/what-are-metaclasses-in-python>

## 11.2 Conventions and patterns

- Subscript at the end of a variable name means we are looking at a coefficient or other computed quantity of an estimator.
- Saving a model using Joblib:

```
1 import joblib
2 # saving; the filetype extension is only for the user/reader
3 joblib.dump(pipe, 'filename.joblib')
4
5 # loading
6 loaded = joblib.load('filename.joblib')
```

Make sure you use the same environment (module version etc.) or the model may change.

- Pipelines are DataFrame-friendly<sup>1</sup>, if the component transformers are as well. If they are not, we can wrap the transformer so that it returns a DataFrame. For instance, `StandardScaler` which normally returns a `np.array`:

```
1 class DFStandardScaler(TransformerMixin):
2     def __init__(self):
3         self.ss = None
4     def fit(self, X, y=None):
5         self.ss = StandardScaler().fit(X)
6         return self
7     def transform(self, X):
8         Xss = self.ss.transform(X)
9         Xscaled = pd.DataFrame(Xss, index=X.index, columns=X.columns)
10        return Xscaled
```

- Pipelines compose, e.g.:

```
1 pipeline = Pipeline([
2     ('features', DFFeatureUnion([
3         ('categoricals', Pipeline([
4             ('extract', ColumnExtractor(CAT_FEATS)),
5             ('dummy', DummyTransformer())
6         ])),
7     ('numerics', Pipeline([
8         ('extract', ColumnExtractor(NUM_FEATS)),
9         ('zero_fill', ZeroFillTransformer()),
10        ('log', Log1pTransformer())
11    ]))
12 ])),
```

---

<sup>1</sup>This info is from [https://www.youtube.com/watch?v=BFaadIqWlAg&list=PLZERW\\_Obpmv\\_t55kNFret-E0hInKeswWF&index=26](https://www.youtube.com/watch?v=BFaadIqWlAg&list=PLZERW_Obpmv_t55kNFret-E0hInKeswWF&index=26), with Github repo at <https://github.com/jem1031/pandas-pipelines-custom-transformers>.

```

13         ('scale', DFStandardScaler())
14     ])

```

### 11.3 Custom Scikit-learn classes

From [https://www.youtube.com/watch?v=WGirN6zBJ4s&list=PLzERW\\_Obpmv\\_t55kNFRet-E0h1nKeswWF&index=1](https://www.youtube.com/watch?v=WGirN6zBJ4s&list=PLzERW_Obpmv_t55kNFRet-E0h1nKeswWF&index=1). There are `Estimator`, `Predictor`, `Transformer`, and `Model` classes. An `Estimator` must implement

- `.fit(X,y)`, fitting the estimator to `X` and `y`
- `.get_params()` return the parameters of the estimator
- `.set_params(**params)` change the parameters of the estimator (e.g. for copying)

Sklearn `Predictor` needs

- `.predict(X)`

Sklearn `Transformer` needs

- `.transform(X, y=None)`

Sklearn `Model` needs

- `.score(X,y)`

#### 11.3.1 Inheritance and Mixins

- You can implement `.get_params()` and `.set_params()` by inheriting from `BaseEstimator`
- There is also `base.TransformerMixin`, `base.RegressorMixin`, `base.ClassifierMixin`, `base.ClusterMixin`, `feature_selection.SelectorMixin`,...

#### 11.3.2 Example 1: simple scaler

```

1 import numpy as np
2 from sklearn.base import BaseEstimator, TransformerMixin
3
4 class Standardizer(BaseEstimator, TransformerMixin):
5
6     def __init__(self, mean_after_transform = 0):
7         self.mean_after_transform = mean_after_transform
8
9     def fit(self, X, y=None):
10         self.mean_ = np.mean(X, axis=0) # columnwise mean
11         self.std_ = np.std(X, axis=0) # columnwise std
12         return self
13
14     def transform(self, X):
15         return (X-self.mean_) / self.std_ + self.mean_after_transform

```

### 11.3.3 Example 2: Regression

Basic regressor that just predicts using the mean or median, using `RegressorMixin` (A regressor is a type of `Model`):

```
1 import numpy as np
2 class MyDummyRegression(BaseEstimator):
3
4     def __init__(self, use_median=False):
5         self.use_median = use_median
6
7     def fit(self, X, y):
8         if self.use_median:
9             self.value_ = np.median(y)
10        else:
11            self.value_ = np.mean(y)
12        return self
13
14    def predict(self, X):
15        out = np.empty (len(X))
16        out.fill(self.value_)
17        return out
```

The mixin gives us a `.score` for free.